

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
df = pd.read_csv("train_dataset.csv")
df.head()
```

	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)	hearing(left)	hearing(right)	systolic	relaxation	...	HDL	LDL	hemoglobin	Urine protein	crea
0	35	170	85	97.0	0.9	0.9	1	1	118	78	...	70	142	19.8	1	
1	20	175	110	110.0	0.7	0.9	1	1	119	79	...	71	114	15.9	1	
2	45	155	65	86.0	0.9	0.9	1	1	110	80	...	57	112	13.7	3	
3	45	165	80	94.0	0.8	0.7	1	1	158	88	...	46	91	16.9	1	
4	20	165	60	81.0	1.5	0.1	1	1	109	64	...	47	92	14.9	1	

5 rows × 23 columns



df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 38984 entries, 0 to 38983
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    38984 non-null  int64
1   height(cm)             38984 non-null  int64
2   weight(kg)              38984 non-null  int64
3   waist(cm)               38984 non-null  float64
4   eyesight(left)          38984 non-null  float64
5   eyesight(right)         38984 non-null  float64
6   hearing(left)           38984 non-null  int64
7   hearing(right)          38984 non-null  int64
8   systolic                38984 non-null  int64
9   relaxation              38984 non-null  int64
10  fasting blood sugar     38984 non-null  int64
11  Cholesterol              38984 non-null  int64
12  triglyceride             38984 non-null  int64
13  HDL                     38984 non-null  int64
14  LDL                     38984 non-null  int64
15  hemoglobin              38984 non-null  float64
16  Urine protein            38984 non-null  int64
17  serum creatinine        38984 non-null  float64
18  AST                     38984 non-null  int64
19  ALT                     38984 non-null  int64
20  Gtp                     38984 non-null  int64
21  dental caries            38984 non-null  int64
22  smoking                 38984 non-null  int64
dtypes: float64(5), int64(18)
memory usage: 6.8 MB
```

```
for i in df.columns:
    print(i + " : ",format(len(df[i].value_counts())))
```

```
age : 14
height(cm) : 13
weight(kg) : 22
waist(cm) : 545
eyesight(left) : 19
eyesight(right) : 17
hearing(left) : 2
hearing(right) : 2
systolic : 125
relaxation : 94
fasting blood sugar : 258
Cholesterol : 279
triglyceride : 389
HDL : 122
LDL : 286
hemoglobin : 143
Urine protein : 6
serum creatinine : 34
AST : 195
ALT : 230
Gtp : 439
dental caries : 2
smoking : 2
```

```
# Check for null values in each column
```

```
null_columns = df.columns[df.isnull().any()]
```

```
# Print the name of columns with null values
```

```
if len(null_columns) > 0:
    print("The following columns contain null values:")
    for col in null_columns:
        print(col)
else:
    print("No null values found in any column.")
```

```
No null values found in any column.
```

```
# plot the graphs of Categorical Variables
```

```
l1 = ['hearing(left)', 'hearing(right)', 'Urine protein', 'dental caries', 'smoking']  
for i in l1:  
    print(i+" : \n", format(df[i].value_counts()))  
    print('\n')
```

```
hearing(left) :  
1      37995  
2       989  
Name: hearing(left), dtype: int64
```

```
hearing(right) :  
1      37963  
2      1021  
Name: hearing(right), dtype: int64
```


```
Urine protein :  
1      36836  
2      1236  
3       667  
4       182  
5        58  
6         5  
Name: Urine protein, dtype: int64
```

```
dental caries :  
0      30625  
1       8359  
Name: dental caries, dtype: int64
```

```
smoking :  
0      24666  
1      14318  
Name: smoking, dtype: int64
```

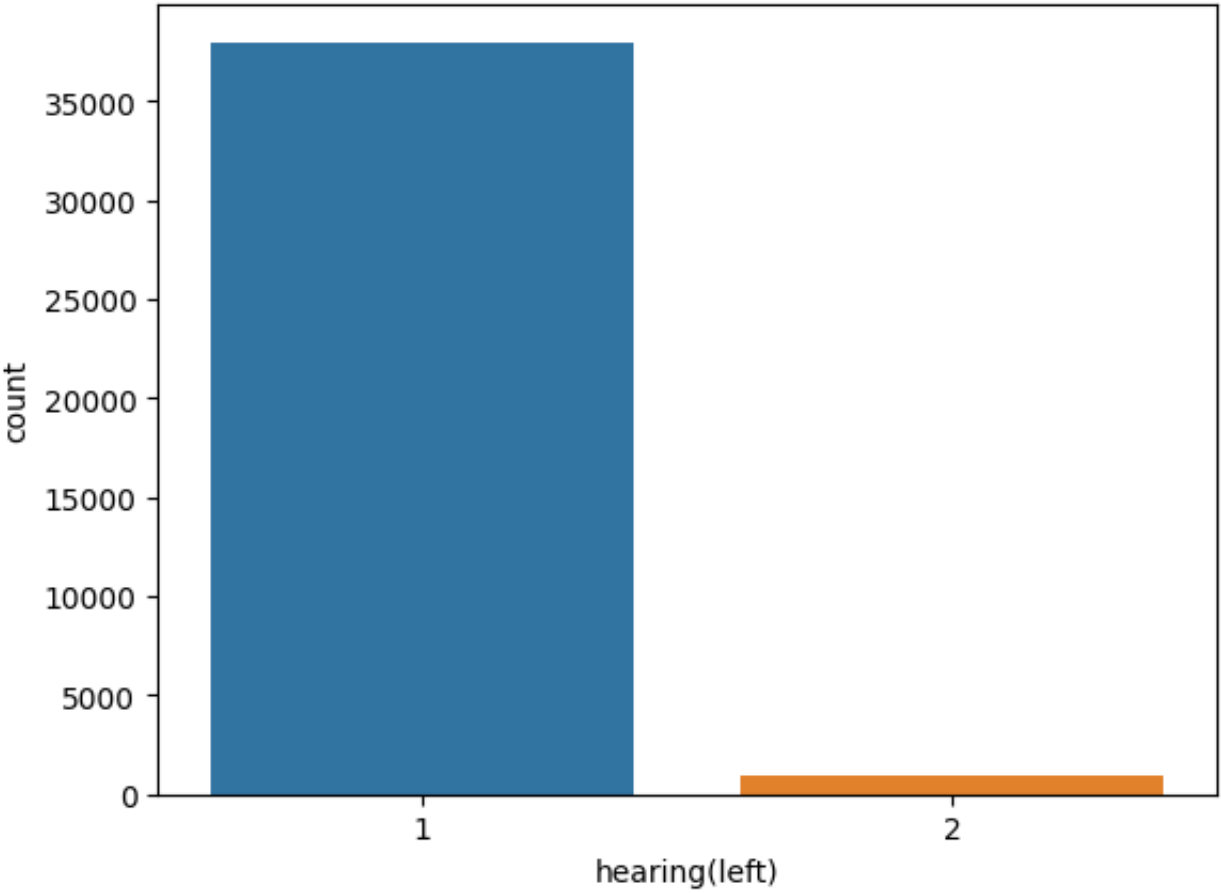
```
sns.countplot(x="smoking", data=df)

plt.show()
```




```
sns.countplot(x="hearing(left)", data=df)

plt.show()
```



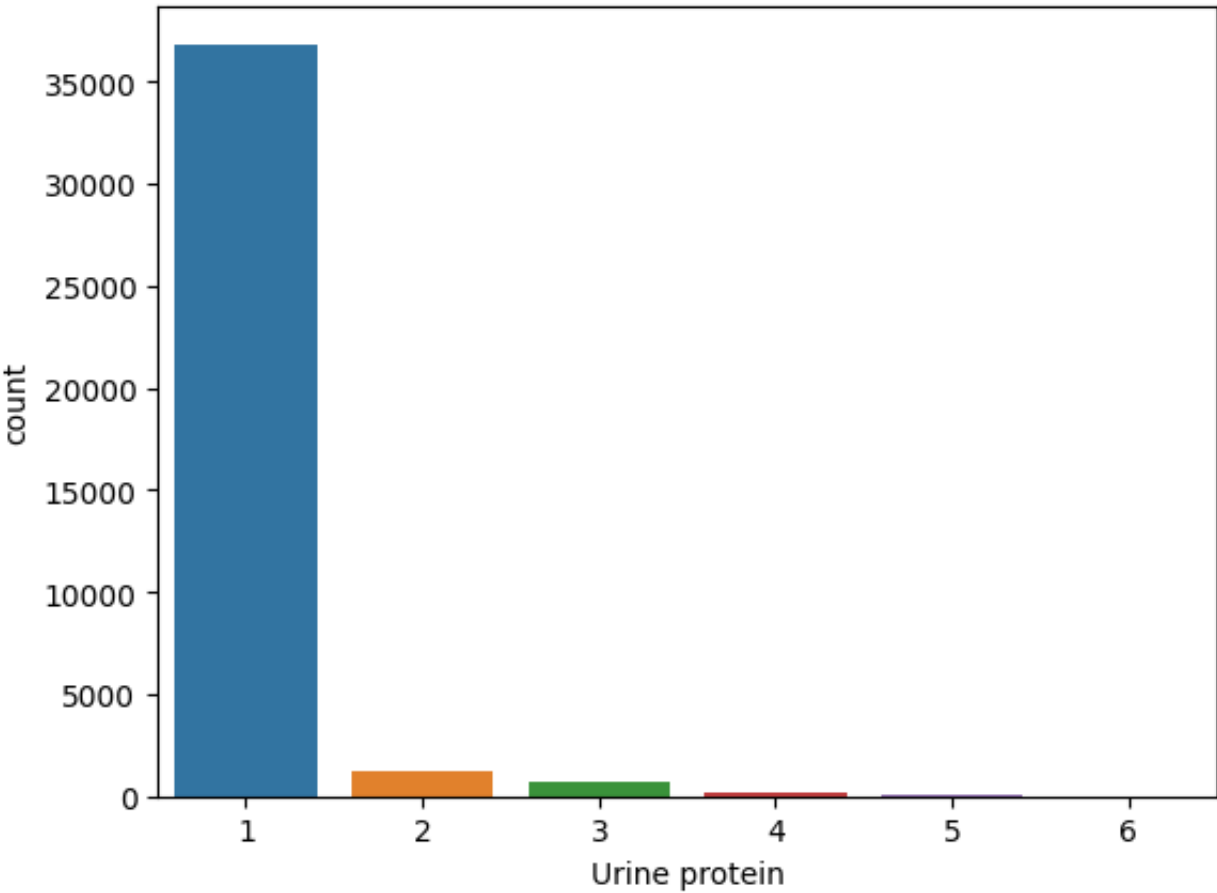
```
sns.countplot(x="hearing(right)", data=df)

plt.show()
```



```
sns.countplot(x="Urine protein", data=df)

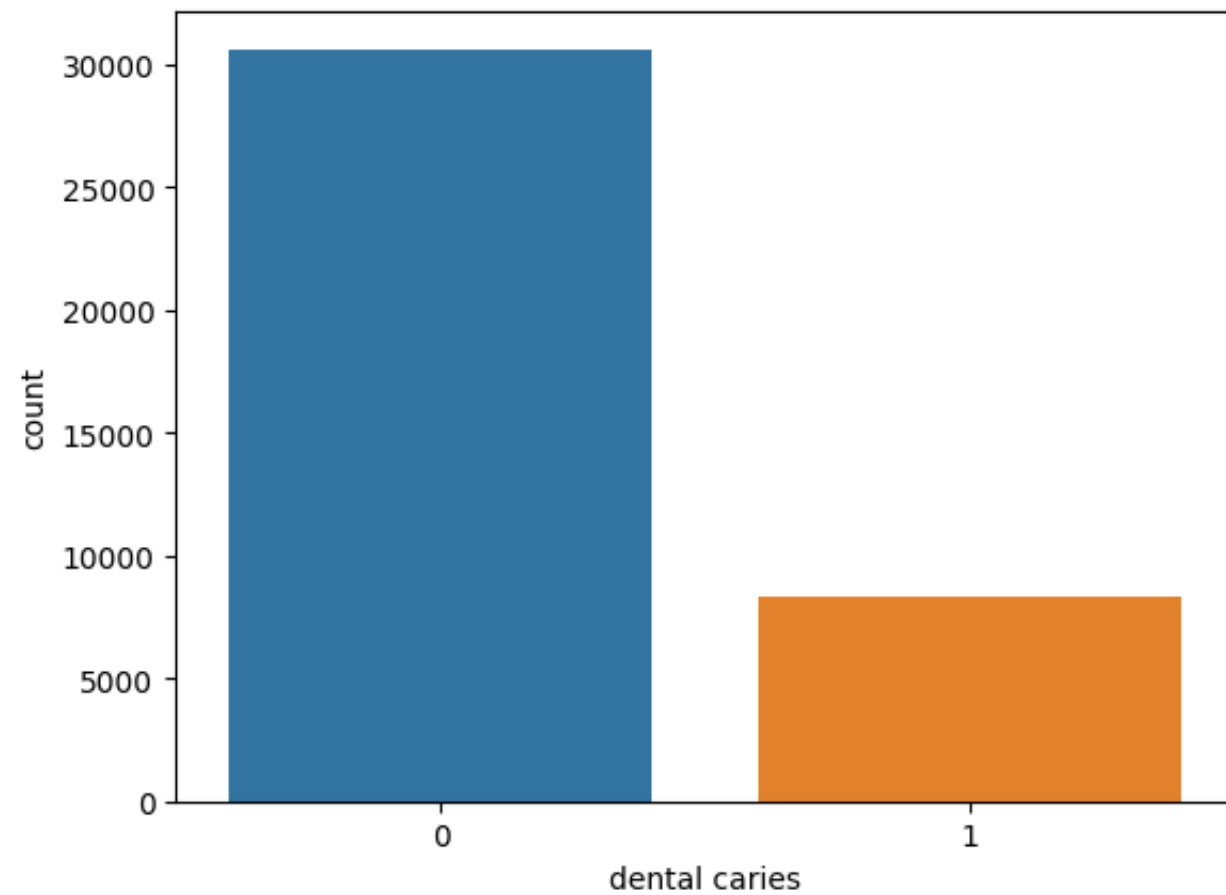
plt.show()
```





```
sns.countplot(x="dental caries", data=df)

plt.show()
```



```
from scipy.stats import chi2_contingency

# Create a contingency table
contingency_table = pd.crosstab(df['smoking'], df['dental caries'])
contingency_table
# Perform chi-square test
stat, p, dof, expected = chi2_contingency(contingency_table)

# interpret p-value
alpha = 0.05
print("Chi-Square test result for Smoking and Dental caries")
print("p value is " + str(p))
if p <= alpha:
    print('Corelated (reject H0)')
else:
    print('Independent (H0 holds true)')

# Create a contingency table
contingency_table = pd.crosstab(df['smoking'], df['Urine protein'])
```

```
contingency_table
# Perform chi-square test
stat, p, dof, expected = chi2_contingency(contingency_table)

# interpret p-value
alpha = 0.05
print("Chi-Square test result for Smoking and Urine protein")
print("p value is " + str(p))
if p <= alpha:
    print('Corelated (reject H0)')
else:
    print('Independent (H0 holds true)')

# Create a contingency table
contingency_table = pd.crosstab(df['smoking'], df['hearing(right)'])
contingency_table
# Perform chi-square test
stat, p, dof, expected = chi2_contingency(contingency_table)

# interpret p-value
alpha = 0.05
print("Chi-Square test result for Smoking and hearing(right)")
print("p value is " + str(p))
if p <= alpha:
    print('Corelated (reject H0)')
else:
    print('Independent (H0 holds true)')


# Create a contingency table
contingency_table = pd.crosstab(df['smoking'], df['hearing(left)'])
contingency_table
# Perform chi-square test
stat, p, dof, expected = chi2_contingency(contingency_table)

# interpret p-value
alpha = 0.05
print("Chi-Square test result for Smoking and hearing(left)")
print("p value is " + str(p))
if p <= alpha:
    print('Corelated (reject H0)')
else:
    print('Independent (H0 holds true)')
```

Chi-Square test result for Smoking and Dental caries  
p value is 4.809679216057804e-100  
Corelated (reject H0)  
Chi-Square test result for Smoking and Urine protein  
p value is 0.049511278397280714  
Corelated (reject H0)  
Chi-Square test result for Smoking and hearing(right)  
p value is 0.0002020239535809226  
Corelated (reject H0)  
Chi-Square test result for Smoking and hearing(left)  
p value is 1.5214418715754427e-05  
Corelated (reject H0)

# Plot the graphs for Numerical Variables

```
fig,axs = plt.subplots(1,6,figsize=(25,5))
axs[0].hist(df['age'])
axs[0].set_title('Age')
axs[1].hist(df['height(cm)'], bins=30, alpha=0.5)
axs[1].set_title('Height(cm)')
axs[2].hist(df['weight(kg)'], bins=30, alpha=0.5)
axs[2].set_title('Weight(kg)')
axs[3].hist(df['waist(cm)'], bins=30, alpha=0.5)
axs[3].set_title('Waist(cm)')
axs[4].hist(df['eyesight(left)'], bins=30, alpha=0.5)
axs[4].set_title('Eyesight(left)')
axs[5].hist(df['eyesight(right)'], bins=30, alpha=0.5)
axs[5].set_title('Eyesight(right)')
```




```


fig,axs = plt.subplots(1,6,figsize=(25,5))
axs[0].hist(df['relaxation'], bins=30, alpha=0.5)
axs[0].set_title('Relaxation')
axs[1].hist(df['fasting blood sugar'], bins=30, alpha=0.5)
axs[1].set_title('Fasting blood sugar')
axs[2].hist(df['Cholesterol'], bins=30, alpha=0.5)
axs[2].set_title('Cholesterol')
axs[3].hist(df['triglyceride'], bins=30, alpha=0.5)
axs[3].set_title('Triglyceride')
axs[4].hist(df['HDL'], bins=30, alpha=0.5)
axs[4].set_title('HDL')
axs[5].hist(df['systolic'], bins=30, alpha=0.5)
axs[5].set_title('systolic')

```

Text(0.5, 1.0, 'systolic')



```
fig,axs = plt.subplots(1,6,figsize=(25,5))
axs[0].hist(df['LDL'], bins=30, alpha=0.5)
axs[0].set_title('LDL')
axs[1].hist(df['hemoglobin'], bins=30, alpha=0.5)
axs[1].set_title('Hemoglobin')
axs[2].hist(df['serum creatinine'], bins=30, alpha=0.5)
axs[2].set_title('Serum creatinine')
axs[3].hist(df['AST'], bins=30, alpha=0.5)
axs[3].set_title('AST')
axs[4].hist(df['ALT'], bins=30, alpha=0.5)
axs[4].set_title('ALT')
axs[5].hist(df['Gtp'], bins=30, alpha=0.5)
axs[5].set_title('Gtp')
```



```
# Remove the outliers
```

```
df_removed_outlier = df[
(df['eyesight(left)'] < 2) &
(df['eyesight(right)'] < 2) &
(df['fasting blood sugar'] < 250) &
(df['triglyceride'] < 400) &
(df['HDL'] < 125) &
(df['LDL'] < 250) &
(df['serum creatinine'] < 2) &
(df['AST'] < 100) &
(df['ALT'] < 200) &
(df['Gtp'] < 300)]

print("Original dataset length: ", len(df))
print("New dataset: ", len(df_removed_outlier))
print("Difference: ", len(df) - len(df_removed_outlier))
```

```
Original dataset length: 38984
New dataset: 37813
Difference: 1171
```

```
# Remove the duplicate rows
```

```
print("Number of duplicate rows: ", df_removed_outlier[df_removed_outlier.duplicated() == True].shape[0])
print(f"\nRows in original Dataframe: {df_removed_outlier.shape[0]}")

new_df = df_removed_outlier.drop_duplicates()
print(f"Dataframe rows after removing duplicates: {new_df.shape[0]}")
```


```
Number of duplicate rows: 5362
```

```
Rows in original Dataframe: 37813
Dataframe rows after removing duplicates: 32451
```

```
# plot the graph after removing outliers and duplicate rows
```

```
fig,axs = plt.subplots(1,6,figsize=(25,5))
axs[0].hist(new_df['age'])
axs[0].set_title('Age')
axs[1].hist(new_df['height(cm)'], bins=30, alpha=0.5)
axs[1].set_title('Height(cm)')
axs[2].hist(new_df['weight(kg)'], bins=30, alpha=0.5)
axs[2].set_title('Weight(kg)')
axs[3].hist(new_df['waist(cm)'], bins=30, alpha=0.5)
axs[3].set_title('Waist(cm)')
axs[4].hist(new_df['eyesight(left)'], bins=30, alpha=0.5)
axs[4].set_title('Eyesight(left)')
axs[5].hist(new_df['eyesight(right)'], bins=30, alpha=0.5)
axs[5].set_title('Eyesight(right)')
```

```
Text(0.5, 1.0, 'Eyesight(right)')
```




```

fig,axs = plt.subplots(1,6,figsize=(25,5))
axs[0].hist(new_df['relaxation'], bins=30, alpha=0.5)
axs[0].set_title('Relaxation')
axs[1].hist(new_df['fasting blood sugar'], bins=30, alpha=0.5)
axs[1].set_title('Fasting blood sugar')
axs[2].hist(new_df['Cholesterol'], bins=30, alpha=0.5)
axs[2].set_title('Cholesterol')
axs[3].hist(new_df['triglyceride'], bins=30, alpha=0.5)
axs[3].set_title('Triglyceride')
axs[4].hist(new_df['HDL'], bins=30, alpha=0.5)
axs[4].set_title('HDL')
axs[5].hist(new_df['systolic'], bins=30, alpha=0.5)
axs[5].set_title('systolic')


```

Text(0.5, 1.0, 'systolic')






```
fig,axs = plt.subplots(1,6,figsize=(25,5))
axs[0].hist(new_df['LDL'], bins=30, alpha=0.5)
axs[0].set_title('LDL')
axs[1].hist(new_df['hemoglobin'], bins=30, alpha=0.5)
axs[1].set_title('Hemoglobin')
axs[2].hist(new_df['serum creatinine'], bins=30, alpha=0.5)
axs[2].set_title('Serum creatinine')
axs[3].hist(new_df['AST'], bins=30, alpha=0.5)
axs[3].set_title('AST')
axs[4].hist(new_df['ALT'], bins=30, alpha=0.5)
axs[4].set_title('ALT')
axs[5].hist(new_df['Gtp'], bins=30, alpha=0.5)
axs[5].set_title('Gtp')
```



```
# Correlation between Numerical variables
```

```
cols = ["age", "height(cm)", "weight(kg)", "waist(cm)", "eyesight(left)", "eyesight(right)", "relaxation",
        "fasting blood sugar", "Cholesterol", "triglyceride", "HDL", "systolic", "LDL", "hemoglobin", "serum creatinine", "AST", "ALT", "Gtp"]
subset = new_df[cols]
```

```
plt.figure(figsize=(20, 10))
plt.title("Correlation matrix for numeric variables")
sns.heatmap(subset.corr(), annot=True, cmap="coolwarm")
plt.show()
```



```
# Dimension Reduction
cols_to_drop = ['waist(cm)', 'systolic', 'Cholesterol', 'AST']

new_df = new_df.drop(columns=cols_to_drop)

new_df.head()
```


	age	height(cm)	weight(kg)	eyesight(left)	eyesight(right)	hearing(left)	hearing(right)	relaxation	fasting blood sugar	triglyceride	HDL	LDL	hemoglobin	Urine protein	s creati
0	35	170	85	0.9	0.9	1	1	78	97	153	70	142	19.8	1	
1	20	175	110	0.7	0.9	1	1	79	88	128	71	114	15.9	1	
3	45	165	80	0.8	0.7	1	1	88	249	366	46	91	16.9	1	
4	20	165	60	1.5	0.1	1	1	64	100	200	47	92	14.9	1	
5	60	160	50	1.0	0.9	2	2	75	114	74	98	64	13.9	1	



```
# Heatmap after dimension reduction

cols = ["age", "height(cm)", "weight(kg)", "eyesight(left)", "eyesight(right)", "relaxation",
        "fasting blood sugar", "triglyceride", "HDL", "LDL", "hemoglobin", "serum creatinine", "ALT", "Gtp"]
subset = new_df[cols]

plt.figure(figsize=(20, 10))
plt.title("Correlation matrix for numeric variables")
sns.heatmap(subset.corr(), annot=True, cmap="coolwarm")
plt.show()
```



```
# Point Biserial test between numerical variables and target categorial variable
```

```
from scipy.stats import pointbiserialr

# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["age"]

# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and age")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')

# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["height(cm)"]

# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and height(cm)")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')

# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["weight(kg)"]

# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and weight(kg)")
print("Point-biserial correlation coefficient:", r_pb)
```

```
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')

# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["eyesight(left)"]

# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and eyesight(left)")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')

# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["eyesight(right)"]

# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and eyesight(right)")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')

# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["relaxation"]
```

```
# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and relaxation")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')

# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["fasting blood sugar"]

# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and fasting blood sugar")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')

# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["triglyceride"]

# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and triglyceride")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')
```

```
# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["HDL"]

# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and HDL")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')
```

```
# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["LDL"]

# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and LDL")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')
```

```
# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["hemoglobin"]

# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
```



```
print("Point-biserial correlation result for Smoking and hemoglobin")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')

# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["serum creatinine"]

# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and serum creatinine")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')

# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["ALT"]

# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and ALT")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')

# Select the binary and continuous variables
binary_var = new_df ["smoking"]
continuous_var = new_df ["Gtp"]
```

```
# Calculate the point-biserial correlation coefficient
r_pb, p_value = pointbiserialr(binary_var, continuous_var)
alpha=0.05
print("Point-biserial correlation result for Smoking and Gtp")
print("Point-biserial correlation coefficient:", r_pb)
print("P-value:", p_value)
if p <= alpha:
    print('Correlated (reject H0)')
else:
    print('Independent (H0 holds true)')
```

```
Point-biserial correlation result for Smoking and age
Point-biserial correlation coefficient: -0.16932003540717697
P-value: 2.851830580380366e-207
Correlated (reject H0)
Point-biserial correlation result for Smoking and height(cm)
Point-biserial correlation coefficient: 0.3957231677997295
P-value: 0.0
Correlated (reject H0)
Point-biserial correlation result for Smoking and weight(kg)
Point-biserial correlation coefficient: 0.3030128002908918
P-value: 0.0
Correlated (reject H0)
Point-biserial correlation result for Smoking and eyesight(left)
Point-biserial correlation coefficient: 0.09535593715018403
P-value: 2.009153736703375e-66
Correlated (reject H0)
Point-biserial correlation result for Smoking and eyesight(right)
Point-biserial correlation coefficient: 0.10408779508383485
P-value: 7.423698239432036e-79
Correlated (reject H0)
Point-biserial correlation result for Smoking and relaxation
Point-biserial correlation coefficient: 0.10271846702778656
P-value: 7.821518695686961e-77
Correlated (reject H0)
Point-biserial correlation result for Smoking and fasting blood sugar
Point-biserial correlation coefficient: 0.09499050920611098
P-value: 6.298966635265718e-66
Correlated (reject H0)
Point-biserial correlation result for Smoking and triglyceride
Point-biserial correlation coefficient: 0.2489267522197282
P-value: 0.0
Correlated (reject H0)
Point-biserial correlation result for Smoking and HDL
Point-biserial correlation coefficient: -0.18321969058712761
P-value: 6.083515653313439e-243
```

```
Correlated (reject H0)
Point-biserial correlation result for Smoking and LDL
Point-biserial correlation coefficient: -0.05684539705552884
P-value: 1.2067548862549454e-24
Correlated (reject H0)
Point-biserial correlation result for Smoking and hemoglobin
Point-biserial correlation coefficient: 0.4008163424929112
P-value: 0.0
Correlated (reject H0)
Point-biserial correlation result for Smoking and serum creatinine
Point-biserial correlation coefficient: 0.25058092348469024
P-value: 0.0
Correlated (reject H0)
Point-biserial correlation result for Smoking and ALT
Point-biserial correlation coefficient: 0.15861909007699465
P-value: 7.840452173195513e-182
Correlated (reject H0)
Point-biserial correlation result for Smoking and Gtp
Point-biserial correlation coefficient: 0.2867833165025298
P-value: 0.0
Correlated (reject H0)
```

```
# Data Partitioning
```

```
new_df['smoking'].value_counts()
```

```
0    20701
1     11750
Name: smoking, dtype: int64
```

```
X = new_df.drop('smoking',axis=1)
y=new_df['smoking']
```

```
from imblearn.over_sampling import SMOTE
```

```
smote=SMOTE()
```

```
X_new,y_new = smote.fit_resample(X,y)
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X_new, y_new, test_size=0.2, random_state=25)
```

## #Logistic Regression

```
from sklearn.linear_model import LogisticRegression
```

```
model = LogisticRegression(random_state=250)
model.fit(X_train, y_train)
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

▼ LogisticRegression

```
LogisticRegression(random_state=250)
```

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import f1_score
```

```
# Accuracy
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

```
# calculate the F1 score
f1 = f1_score(y_test, y_pred)
```

```
print('F1 score:', f1)
```

```
# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)
```

```
# Get TP, FP, TN, FN
tn, fp, fn, tp = cm.ravel()
```

```
# Calculate sensitivity and specificity
sensitivity = tp / (tp + fn)
specificity = tn / (tn + fp)
```


```
# Print the results
print("Sensitivity:", sensitivity)
print("Specificity:", specificity)
```

```
# Visualize confusion matrix as heatmap
sns.heatmap(cm, annot=True, fmt='g')
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.title('Confusion matrix')
plt.show()


# Compute ROC curve and AUC score
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
auc = roc_auc_score(y_test, y_pred)
print('AUC score:', auc)

# Plot ROC curve
plt.plot(fpr, tpr, color='darkorange', label='ROC curve')
plt.plot([0, 1], [0, 1], color='navy', linestyle='--', label='Random guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve (AUC = {:.2f})'.format(auc))
plt.show()
```

Accuracy: 0.7131988890230648  
F1 score: 0.717026093172882  
Sensitivity: 0.7367776689520078  
Specificity: 0.6902549440076245



AUC score: 0.7135163064798161



```
#Decision tree classifier

from sklearn.tree import DecisionTreeClassifier

model = DecisionTreeClassifier(min_samples_split=5, random_state=42)
model.fit(X_train, y_train)
```

▼ DecisionTreeClassifier

DecisionTreeClassifier(min\_samples\_split=5, random\_state=42)

```
# Accuracy
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# calculate the F1 score
f1 = f1_score(y_test, y_pred)

print('F1 score:', f1)
```

```
# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Get TP, FP, TN, FN
tn, fp, fn, tp = cm.ravel()

# Calculate sensitivity and specificity
sensitivity = tp / (tp + fn)
specificity = tn / (tn + fp)

# Print the results
print("Sensitivity:", sensitivity)
print("Specificity:", specificity)

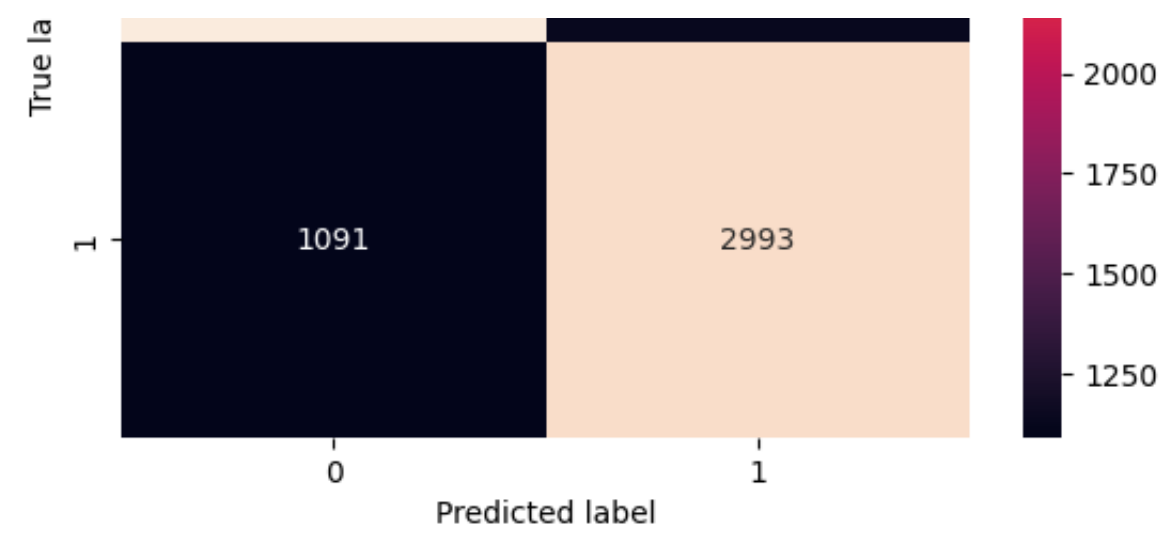
# Visualize confusion matrix as heatmap
sns.heatmap(cm, annot=True, fmt='g')
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.title('Confusion matrix')
plt.show()

# Compute ROC curve and AUC score
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
auc = roc_auc_score(y_test, y_pred)
print('AUC score:', auc)


# Plot ROC curve
plt.plot(fpr, tpr, color='darkorange', label='ROC curve')
plt.plot([0, 1], [0, 1], color='navy', linestyle='--', label='Random guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve (AUC = {:.2f})'.format(auc))
plt.show()
```

```
Accuracy: 0.7317956768506219
F1 score: 0.7293773607895698
Sensitivity: 0.7328599412340843
Specificity: 0.7307600667143197
```





AUC score: 0.731810003974202





```
# Random Forest Classifier
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
```

```
▼ RandomForestClassifier
RandomForestClassifier(random_state=42)
```

```
# Accuracy
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

```
# calculate the F1 score
f1 = f1_score(y_test, y_pred)
```

```
print('F1 score:', f1)
```

```
# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)
```

```
# Get TP, FP, TN, FN
tn, fp, fn, tp = cm.ravel()
```

```
# Calculate sensitivity and specificity
sensitivity = tp / (tp + fn)
specificity = tn / (tn + fp)
```


```
# Print the results
print("Sensitivity:", sensitivity)
print("Specificity:", specificity)
```

```
# Visualize confusion matrix as heatmap
sns.heatmap(cm, annot=True, fmt='g')
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.title('Confusion matrix')
plt.show()
```


```
# Compute ROC curve and AUC score
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
auc = roc_auc_score(y_test, y_pred)
print('AUC score:', auc)
```


```
# Plot ROC curve
plt.plot(fpr, tpr, color='darkorange', label='ROC curve')
plt.plot([0, 1], [0, 1], color='navy', linestyle='--', label='Random guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve (AUC = {:.2f})'.format(auc))
plt.show()
```

Accuracy: 0.8085979954111823  
F1 score: 0.81537565521258  
Sensitivity: 0.8570029382957884  
Specificity: 0.7614963068858709



AUC score: 0.8092496225908297





```
# KNN classifier
```

```
from sklearn.neighbors import KNeighborsClassifier
# Create a KNN classifier with k=3
model = KNeighborsClassifier(n_neighbors=3, weights='distance', metric='euclidean')
```

```
# Fit the classifier to the training data
model.fit(X_train, y_train)
```

```
▼ KNeighborsClassifier
KNeighborsClassifier(metric='euclidean', n_neighbors=3, weights='distance')
```

```
# Accuracy
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

```
# calculate the F1 score
f1 = f1_score(y_test, y_pred)
```

```
print('F1 score:', f1)
```

```
# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)
```

```
# Get TP, FP, TN, FN
tn, fp, fn, tp = cm.ravel()
```

```
# Calculate sensitivity and specificity
sensitivity = tp / (tp + fn)
specificity = tn / (tn + fp)
```

```
# Print the results
```

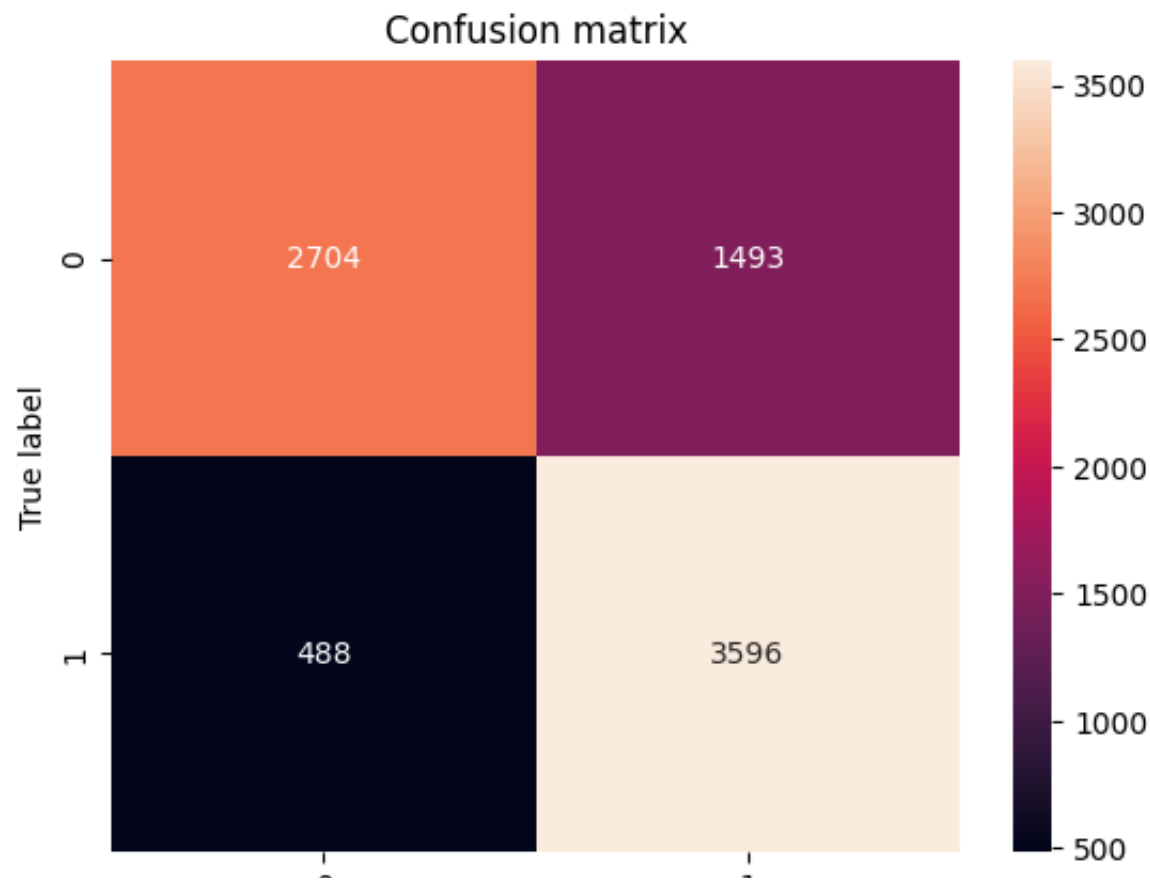
```
print("Sensitivity:", sensitivity)
print("Specificity:", specificity)

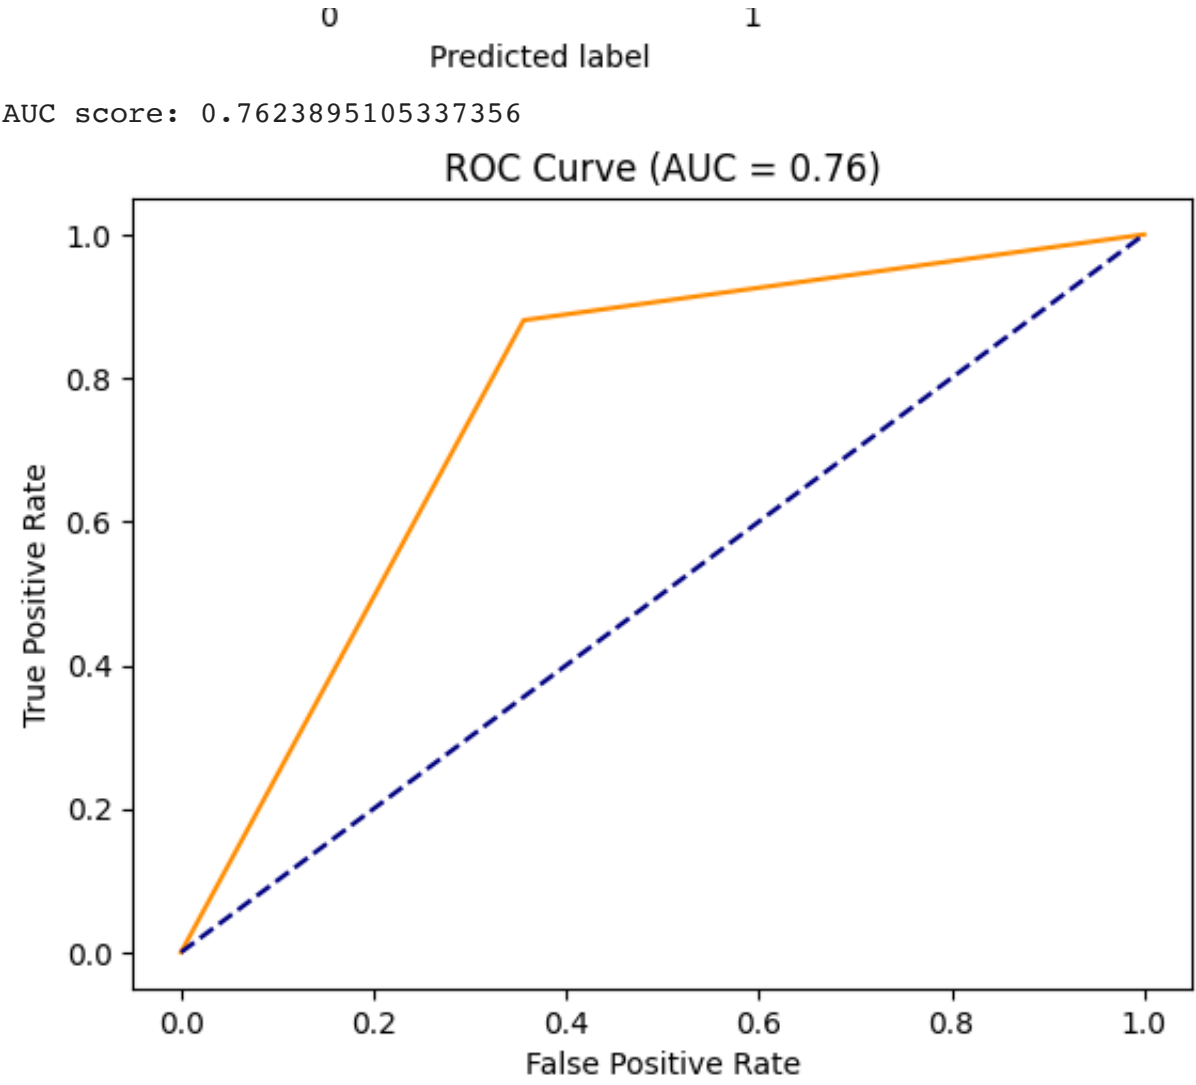
# Visualize confusion matrix as heatmap
sns.heatmap(cm, annot=True, fmt='g')
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.title('Confusion matrix')
plt.show()

# Compute ROC curve and AUC score
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
auc = roc_auc_score(y_test, y_pred)
print('AUC score:', auc)

# Plot ROC curve
plt.plot(fpr, tpr, color='darkorange', label='ROC curve')
plt.plot([0, 1], [0, 1], color='navy', linestyle='--', label='Random guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve (AUC = {:.2f})'.format(auc))
plt.show()
```

Accuracy: 0.760777683854607  
F1 score: 0.7840401177368364  
Sensitivity: 0.8805093046033301  
Specificity: 0.6442697164641411






```
# data
x = ['Logistic Regression', 'Decision tree', 'Random Forest', 'KNN classifier']
y = [0.712, 0.740, 0.809, 0.764]

# create a bar chart
plt.bar(x, y)

# add labels and title
plt.xlabel('Model')
plt.ylabel('Accuracy')
plt.title('Comparision of Models')
plt.xticks(rotation=45)

# show the plot
plt.show()
```



[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s    completed at 4:59 PM

